

## Lab 1: Vector Addition

Due Date: February 6 (Thursday), 2020 at 11:59pm

### 1. Objective

The purpose of this lab is to get you familiar with using the CUDA API by implementing a simple vector addition kernel and its associated setup code.

### 2. Procedure

**Step 1:** Download the Lab 1 materials from blackboard to your home directory on the Karpinski GPU cluster. Unzip the file.

```
unzip lab1-vector-add.zip
```

**Step 2:** Edit the file `lab1-vector-add/main.cu` to implement the following where indicated:

- a) Allocate device memory
- b) Copy host memory to device
- c) Initialize thread block and kernel grid dimensions and invoke CUDA kernel
- d) Copy results from device to host
- e) Free device memory

**Step 3:** Edit the file `lab1-vector-add/kernel.cu` to implement the vector addition kernel code.

**Step 4:** Compile and test your code.

```
cd lab1-vector-add  
  
make  
  
./vecadd          # Uses the default vector size  
./vecadd m        # Uses vectors of size m
```

**Step 5:** Submit your assignment. You only need to submit the following files:

- `main.cu`
- `kernel.cu`

Compress the files and name them after your last name as follows:

```
tar -cvf lab1_<lastname>.tar main.cu kernel.cu
```

Submit the compressed file to blackboard before the deadline.

### 3. Grading:

The full grade of this lab is 100 points, which are broken down as follows.

- Correctness (60 points): both default vector size and user-input vector size will be tested.
  - If code is not functioning correctly, grade will be zero.
- Completeness (30 points): all the missing parts are completed.
- Test on the return values of APIs (10 points): check if the calling of APIs is successful.
  - See the last slide of Lecture 2 for example.

*Note: This is a simple but essential exercise. Please write out the code and do not copy it from other examples or lecture slides. That process is most important.*